

Short version: nothing is “wrong” with your biology, the issue is your numerics. Forward Euler on a logistic-like ODE **does not** guarantee positivity. If your time step is too big relative to the effective growth rate and carrying capacity, the scheme can overshoot and drive a state negative—even if the underlying continuous model can’t.

Let’s unpack that in your specific code.

1. Your update equation *can* go negative

Take one of your state updates, e.g. for leaves:

$$N[t+1] = N[t] + dt * (aN_hat * N[t] * (1 - N[t]/max(kN_hat, 1e-9)))$$

This is Forward Euler applied to the logistic ODE

$$\left[\frac{dN}{dt} = a_N \text{hat}, N \left(1 - \frac{N}{k_N \text{hat}}\right) \right]$$

In continuous time, with $(N(0) \geq 0)$, the solution stays (≥ 0) forever.

In discrete time with Euler, you’re doing:

$$\left[N_{t+1} = N_t + \Delta t, a_N \text{hat}, N_t \left(1 - \frac{N_t}{k_N \text{hat}}\right) \right]$$

Factor out (N_t) :

$$\left[N_{t+1} = N_t \Big[1 + \Delta t, a_N \text{hat} \Big] \left(1 - \frac{N_t}{k_N \text{hat}}\right) \right]$$

So:

- If $(N_t < k_N \text{hat})$, the term in parentheses is positive, so you grow.
- If $(N_t > k_N \text{hat})$, that term is negative, which is how the model “pushes” you back toward the carrying capacity.

But if the negative part is **too large in magnitude**, the bracket can become negative:

$$[1 + \Delta t, a_{N,\text{hat}} \left(1 - \frac{N_t}{k_{N,\text{hat}}}\right) < 0]$$

which implies

$$[\Delta t, a_{N,\text{hat}} \left(\frac{N_t}{k_{N,\text{hat}}} - 1\right) > 1.]$$

When that inequality holds, you get ($N_{t+1} < 0$).

So even without any explicit decay term, **Euler stepping can overshoot below zero** once a state gets above its effective carrying capacity and the time step is too large for the local dynamics.

2. Why it hits leaves / spikelets especially

The part that makes this more likely in your code is the time-varying effective parameters:

```
# Nutrient factors (bounded, nonnegative)
nuW = get_nutrient_factor(...)
nuF = get_nutrient_factor(...)
nuT = get_nutrient_factor(...)
nuR = get_nutrient_factor(...)

# Growth rates
aN_hat = aN           # simpler here, but aN might already be big

# Carrying capacities
kN_hat = kN * (nuT * nuR)**(1/2)
```

Things to notice:

- If **nuT** or **nuR** get very small (e.g., poor temperature / radiation), then $(\nu_T \nu_R)^{1/2}$ can be tiny → **kN_hat becomes very small**.
- If (N_t) hasn't dropped correspondingly, the ratio (N_t / kN_{hat}) can become huge.
- Then the logistic term $(1 - N_t / kN_{\text{hat}})$ is a large negative number, and the Euler step happily takes a huge negative step.

Same logic applies to flower spikelets if their (a) and (k) get into ugly combinations via the nutrient factors.

You already avoid divide-by-zero with:

```
max(kN_hat, 1e-9)
```

but that only prevents NaNs; it does **not** prevent “massive negative derivative times dt”.

So the root cause is:

Effective carrying capacity gets very small or state gets too large, and the Euler step is too big to stay in $[0, k]$.

3. This is a known Euler problem, not you reinventing chaos (on purpose)

Forward Euler applied to stiff or strongly nonlinear systems is notorious for:

- breaking invariants (like non-negativity),
- overshooting,
- and even producing chaotic behavior in otherwise tame continuous systems.

Your model is basically a **time-varying logistic map** embedded in a larger system. Logistic maps are famous for doing weird stuff when parameters are large; Euler is turning your smooth ODE into something closer to that discrete logistic map in regimes with big effective growth.

4. Ways to fix or mitigate it

Here are practical knobs you can turn:

a) Reduce `dt` or substep the integration

Easiest conceptual fix:

- Make `dt` smaller so that ($\Delta t, a_{\text{eff}}(N_t/k_{\text{eff}} - 1)$) stays well below 1 even in the worst case.

If you don't want to change global `dt`, you can substep inside the loop when conditions are harsh:

```
substeps = 4 # or adapt based on aN_hat, kN_hat, etc.
sub_dt = dt / substeps
for _ in range(substeps):
    N[t+1] = N[t] + sub_dt * (aN_hat * N[t] * (1 - N[t]/max(kN_hat, 1e-9)))
    N[t+1] = max(N[t+1], 0.0) # see (c)
    N[t] = N[t+1]
```

You'd do similarly for `h, A, c, P`.

b) Tame the nutrient scaling

If your `get_nutrient_factor` can produce values very close to 0 or very large, the resulting `k_hat` and `a_hat` can explode or collapse.

Common trick: clamp nutrient factors to reasonable ranges, e.g.:

```
nuW = np.clip(get_nutrient_factor(...), 0.1, 2.0)
```

You can also clamp the effective carrying capacities:

```
kN_hat = np.clip(kN_hat, kN_min, kN_max)
```

where `kN_min` is some biologically reasonable minimum (so you never “compress” the carrying capacity to nearly zero).

c) Enforce non-negativity explicitly

Not elegant mathematically, but very common in ecological / biological simulations:

```
h[t+1] = max(h[t+1], 0.0)
A[t+1] = max(A[t+1], 0.0)
N[t+1] = max(N[t+1], 0.0)
c[t+1] = max(c[t+1], 0.0)
P[t+1] = max(P[t+1], 0.0)
```

Or even:

```
if N[t+1] < 0:
    N[t+1] = 0
```

Physically, negative leaves don't exist; biologically you might reinterpret this as “plant lost almost all leaves, state floored at 0.”

For counts (leaves, spikelets), another fancy option is to model them in log-space or with a softplus:

- Evolve a variable (`z`), and define (`N = \exp(z)`) or (`N = \text{softplus}(z)`), which is always positive.
- This requires rewriting the dynamics, though, so it's more invasive.

d) Use a more stable integrator

If performance permits, you can switch from hand-coded Euler to something like Runge–Kutta (RK4) or a semi-implicit method.

For example, you could implement a simple RK4 step for each variable. It still won't *guarantee* positivity, but it's much more stable for the same dt .

Since you're calling this inside a GA cost function, you'd need to balance:

- number of individuals \times number of time steps \times integrator complexity.

But even a low-overhead improvement like Heun's method (a 2-stage Runge–Kutta) can help a lot.

5. How to debug *where* it happens

A quick practical check:

```
for t in range(total_time_steps - 1):
    ...
    N_pred = N[t] + dt * (aN_hat * N[t] * (1 - N[t]/max(kN_hat, 1e-9)))
    if N_pred < 0:
        print("Negative N step at t=", t)
        print("N[t]=", N[t], "aN_hat=", aN_hat, "kN_hat=", kN_hat, "dt=", dt)
    N[t+1] = max(N_pred, 0.0)
```

My bet: you'll catch cases where kN_hat is extremely small or $aN_hat \cdot dt$ effectively huge.

So the punchline:

- The model equations themselves are fine: logistic growth can go down but not below zero in continuous time.
- The **combination of Forward Euler + time-varying tiny carrying capacities / large growth rates + finite dt** lets the discrete approximation overshoot below zero.
- Fix it by shrinking dt or substepping, clamping extremes in nutrient scaling and capacities, and optionally enforcing non-negativity.

You've basically built a little nonlinear dynamical system zoo; what you're seeing is it misbehaving in the discrete-time cage, not in the continuous-time wild.

For speed, I don't want to decrease dt . Would using RK4 instead of Forward Euler help?

Short answer: RK4 will *probably* reduce how often things blow up or go negative for the same dt , but it **does not guarantee** positivity. It's a nicer Euler, not a magic physics cop.

So if your constraint is “I really don't want to shrink dt ”, I'd consider three options, in roughly this order:

1. RK4: what it helps and what it doesn't

Right now a single state does:

```
N[t+1] = N[t] + dt * f(N[t])
# where f(N) = aN_hat * N * (1 - N/kN_hat)
```

Euler is first-order and has a small stability region. RK4 replaces that by:

```
k1 = f(N_t)
k2 = f(N_t + 0.5*dt*k1)
k3 = f(N_t + 0.5*dt*k2)
k4 = f(N_t + dt*k3)
N_{t+1} = N_t + dt*(k1 + 2*k2 + 2*k3 + k4)/6
```

With your code structure, note something nice:

- All the **nutrient factors** and **effective parameters** a^*_hat , k^*_hat depend only on t and the cumulative drivers (WC , RC , etc.), **not** on the plant states.
- That means inside each time step, the RHS is just $(f(h, A, N, c, P))$ with *fixed* a^*_hat , k^*_hat .

So you can:

1. Compute nutrient factors + a^*_hat , k^*_hat **once per step**.
2. Use those in RK4 stages for all 5 state variables.

Cost-wise, that's $\sim 4x$ more multiplications/adds for the states, but **no extra nutrient-factor calls**. For a GA inner loop this might be acceptable.

This will:

- Reduce overshoot,
- Give you better accuracy for the same dt ,
- Shrink the region where you get wild behavior.

But: in extreme cases (tiny k_hat , huge a_hat , large state), RK4 can still drive you negative. Explicit RK methods do not enforce invariants like “state ≥ 0 ”.

So: **yes, RK4 helps stability, no, it doesn't solve the positivity problem in a guaranteed way.**

2. A very nice alternative: exact logistic step (same cost ballpark)

Because each variable follows a logistic-like ODE with *constant* a_{hat} and k_{hat} over the step, you can actually use the **closed-form solution** for each step instead of Euler.

For:

$$[\frac{dy}{dt} = a_y \left(1 - \frac{y}{K}\right)]$$

with constant ($a, K > 0$) over the step, the exact step from (y_n) to (y_{n+1}) over dt is:

$$[y_{n+1} = \frac{K}{1 + \left(\frac{K}{y_n} - 1\right)e^{-a dt}}.]$$

This has two big perks:

- If $(y_n > 0, a \geq 0, K > 0)$, then $(y_{n+1} > 0)$ automatically.
- It's *very* stable; you don't need tiny dt for the logistic part.

You could implement a helper like:

```
def logistic_step(y, a, K, dt, eps=1e-12):  
    K = max(K, eps)  
    y = max(y, eps)  
    exp_term = np.exp(-a * dt)  
    return K / (1.0 + (K / y - 1.0) * exp_term)
```

Then inside your loop:

```
h[t+1] = logistic_step(h[t], ah_hat, kh_hat, dt)  
A[t+1] = logistic_step(A[t], aa_hat, ka_hat, dt)  
N[t+1] = logistic_step(N[t], an_hat, kn_hat, dt)  
c[t+1] = logistic_step(c[t], ac_hat, kc_hat, dt)  
P[t+1] = logistic_step(P[t], ap_hat, kp_hat, dt)
```

This is still per-step work on the same order as Euler + a few `exp` calls, and far cheaper than cutting dt by a factor of, say, 4–10.

If any of your a_{hat} can be negative, you'd want to think a bit about what "logistic" means biologically there, but for your use-case they look like growth rates.

3. Cheap guardrails if you stay with Euler (or RK4)

Regardless of method, you can add safety rails with negligible cost:

- Clamp nutrient factors / capacities so k^* never gets absurdly tiny.
- Enforce non-negativity after stepping:

```
h[t+1] = max(h[t+1], 0.0)
A[t+1] = max(A[t+1], 0.0)
N[t+1] = max(N[t+1], 0.0)
c[t+1] = max(c[t+1], 0.0)
P[t+1] = max(P[t+1], 0.0)
```

For a GA objective, this kind of clipping is often totally acceptable: a schedule that tries to drive plants into physically impossible regimes just gets penalized.

TL;DR for design choice

- **RK4**: better stability and accuracy, same Δt , $\sim 4 \times$ state-update cost, still not positivity-guaranteed.
- **Exact logistic step**: very stable, naturally positivity-preserving (under mild assumptions), a bit more math but similar cost per step.
- **Clamping + Euler**: simplest hack, numerically crude but often absolutely fine for GA.

If your priority stack is “speed > elegance > ironclad positivity”, I’d do:

1. Switch each state update to the exact logistic step;
2. Add a floor at 0 as a final guard;
3. Optionally clamp extreme nutrient factors/carrying capacities.

That way your simulated plants don’t mysteriously develop negative flowers, even when your GA explores weird corners of design space.